

Shard Value Summary (v0.6.2)

What Shard Is:

- A distributed inference network combining browser scouts and verifier nodes.
- OpenAI-compatible API plus overflow routing and SLA controls.

Cost Savings:

- Contribution mode enables compute-for-compute usage without per-token API charges.
- Reference API costs for centralized providers commonly range from 0.002to0.06 per 1K tokens.

Performance (Validated Current Assessment):

- Only validated node-count benchmark points are plotted; missing points are marked pending.
- Phase 3 local 1000-scout: acceptance 0% | error 100% | p95 2300 ms | requests 11717.
- Phase 3 EC2 1000-scout quick: acceptance 0% | error 0.1178% | p95 2.011e+04 ms | requests 5940.
- Phase 3 EC2 100-scout baseline: acceptance 100% | error 9.167% | p95 3.001e+04 ms | requests 1189.

Phase 4 SLA Snapshot (latest report):

- SLA Breaches: 0.0000 vs == 0 over 7d -> PASS
- Weekly p95 (ms): 0.0000 vs <= 3000 -> PASS
- Bootstrap Uptime Ratio: 0.0000 vs >= 0.995 -> FAIL
- Acceptance Rate: 0.0000 vs >= 0.65 -> FAIL

Business Problems Solved:

- Cost control: reduce dependency on per-token third-party billing.
- Scaling resilience: add distributed compute contributors during traffic spikes.
- Ownership: run inference fabric under your own infrastructure and policies.

Participation Incentives:

- Scout contributors donate idle browser compute and receive reciprocal network utility.
- Verifier operators gain reliable overflow capacity for their own workloads.
- No token is required to start contributing in the current model.

Call to Action:

- Scout: open shardnetwork.live and join in under a minute.
- Verifier: run `docker compose shard-daemon` and expose mesh ports.