

Shard Value Summary (v0.6.2)

What Shard Is:

- A distributed inference network combining browser scouts and verifier nodes.
- OpenAI-compatible API plus overflow routing and SLA controls.

Cost Savings:

- Contribution mode enables compute-for-compute usage without per-token API charges.
- Reference API costs for centralized providers commonly range from 0.002to0.06 per 1K tokens.

Performance (Current Assessment):

- Staging benchmark path demonstrates low latency under 1000-scout synthetic drill conditions.
- Throughput and error-rate gates can pass with tuned orchestration.
- Speculative acceptance telemetry is still being improved for production-grade validation.

Business Problems Solved:

- Cost control: reduce dependency on per-token third-party billing.
- Scaling resilience: add distributed compute contributors during traffic spikes.
- Ownership: run inference fabric under your own infrastructure and policies.

Participation Incentives:

- Scout contributors donate idle browser compute and receive reciprocal network utility.
- Verifier operators gain reliable overflow capacity for their own workloads.
- No token is required to start contributing in the current model.

Call to Action:

- Scout: open shardnetwork.live and join in under a minute.
- Verifier: run `docker compose shard-daemon` and expose mesh ports.